

# Abstract

This report details the implementation and evaluation of Recurrent Neural Networks (RNN), Long Short-Term Memory (LSTM), for forecasting daily minimum grass temperature in Hong Kong. Using historical data from the Hong Kong Observatory (1980–2024 for training and January–October 2025 for testing), we preprocessed univariate time-series data, trained models in TensorFlow/Keras, and assessed performance with MAE and RMSE metrics. RNN achieved the best results (MAE=0.96°C, RMSE=1.31°C), outperforming LSTM, with insights into Hong Kong's subtropical climate challenges like seasonal extremes. All code, models, and results, and results are in our public GitHub repository:

[<https://github.com/ben12345qq/Developing-RNN-and-LSTM-Models-for-Hong-Kong-Daily-Minimum-Temperature-Forecasting.git>].

## Methodology

### Data Acquisition and Preprocessing

The dataset (CSV format) covers daily grass minimum temperatures (°C) at HKO headquarters (22.3°N, 114.17°E, 134m elevation) from 1968–2025. We downloaded from <https://data.gov.hk/en-data/dataset/hk-hko-rss-daily-grass-min-temp>.

Preprocessing steps (from 5411Project.py):

Skip 2 metadata rows; strip/ rename bilingual columns (e.g., '年/Year' to 'year').

Coerce invalid values ('\*\*\*' for missing) to NaN; drop footer rows.

Create datetime index; filter complete data ('C' in completeness column).

Interpolate/fill remaining NaNs; subset training (1980–2024: ~16,436 days) and test (2025: ~304 days).

Scale with MinMaxScaler (0–1); create 30-day sequences for next-day prediction.

## 5411Project.py code

```
import pandas as pd

from sklearn.preprocessing import MinMaxScaler

import numpy as np


from tensorflow.keras.layers import Conv1D, MaxPooling1D, Flatten
import matplotlib.pyplot as plt

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import SimpleRNN, LSTM, Dense, Dropout,
Bidirectional
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping
import pickle
import os

os.environ['TF_ENABLE_ONEDNN_OPTS'] = '0' # Suppress oneDNN warnings

# Inspect the file
df = pd.read_csv('.\daily_HKO_GMT_ALL.csv', skiprows=2, encoding='utf-8')
print(df.head(5))

df.columns = [col.strip() for col in df.columns]

# Rename to standard English lowercase
df = df.rename(columns={

    '年/Year': 'year',
```

```
'月/Month': 'month',
'日/Day': 'day',
'數值/Value': 'value',
'數據完整性/data Completeness': 'data_completeness'

})

# Convert to numeric, coerce invalid to NaN (handles '***', footer strings, etc.)

df['year'] = pd.to_numeric(df['year'], errors='coerce')

df['month'] = pd.to_numeric(df['month'], errors='coerce')

df['day'] = pd.to_numeric(df['day'], errors='coerce')

df['value'] = pd.to_numeric(df['value'], errors='coerce')      # Handles '***' in
temperatures

# Drop rows with NaN in date columns (removes footer and invalid dates)

df = df.dropna(subset=['year', 'month', 'day'])

# Now safely cast date components to integers

df['year'] = df['year'].astype(int)

df['month'] = df['month'].astype(int)

df['day'] = df['day'].astype(int)

# Create datetime index

df['Date'] = pd.to_datetime(df[['year', 'month', 'day']], errors='raise')  # Raise if any
issues remain

df = df.set_index('Date')
```

```
# Filter for complete data only and select temperature (excludes incomplete or
missing values)

df = df[df['data_completeness'] == 'C']['value']

# Handle any remaining missing temperatures (though 'C' should mean no misses)

df = df.interpolate(method='linear').ffill().bfill()

# Subsets

train_data = df['1980-01-01':'2024-12-31']

test_data = df['2025-01-01':'2025-10-30']

# Verify

print(f"Train data shape: {train_data.shape}") # Should be ~16436 days (45 years,
accounting for leap years/missing)

print(f"Test data shape: {test_data.shape}") # Should be ~304 days (Jan-Oct
2025)

print(train_data.head())

print(test_data.tail())

# Handle missing values

train_data = train_data.interpolate(method='linear').ffill().bfill() # Interpolate,
then fill edges

# Plot for analysis
```

```

plt.figure(figsize=(12, 6))

plt.plot(train_data)

plt.title('HK Daily Grass Min Temperature (1980-2024)')

plt.xlabel('Date')

plt.ylabel('Temperature (°C)')

plt.show()

# Scale data

scaler = MinMaxScaler(feature_range=(0, 1))

train_scaled = scaler.fit_transform(train_data.values.reshape(-1, 1))

test_scaled = scaler.transform(test_data.values.reshape(-1, 1)) # Use same
scaler

# Create sequences (e.g., window_size=30)

def create_sequences(data, window_size):

    X, y = [], []

    for i in range(len(data) - window_size):

        X.append(data[i:i+window_size])

        y.append(data[i+window_size])

    return np.array(X), np.array(y)

window_size = 30 # Experiment with 30-60

X_train, y_train = create_sequences(train_scaled, window_size)

X_test, y_test = create_sequences(test_scaled, window_size)

```

```
# Reshape for RNN/LSTM (samples, timesteps, features)

X_train = X_train.reshape((X_train.shape[0], X_train.shape[1], 1))

X_test = X_test.reshape((X_test.shape[0], X_test.shape[1], 1))
```

with open('scaler.pkl', 'wb') as f:

```
pickle.dump(scaler, f)

np.save('X_test.npy', X_test)

np.save('y_test.npy', y_test)

test_data.to_pickle('test_data.pkl')
```

with open('window\_size.txt', 'w') as f:

```
f.write(str(window_size))

# RNN Model

rnn_model = Sequential()

rnn_model.add(SimpleRNN(50,           input_shape=(window_size,      1),
                       return_sequences=True))

rnn_model.add(SimpleRNN(50))

rnn_model.add(Dense(1))

rnn_model.compile(optimizer=Adam(learning_rate=0.001), loss='mse')
```

# LSTM Model

```
lstm_model = Sequential()

lstm_model.add(LSTM(50,           input_shape=(window_size,      1),
                     return_sequences=True))
```

```
lstm_model.add(Dropout(0.2))

lstm_model.add(LSTM(50))

lstm_model.add(Dropout(0.2))

lstm_model.add(Dense(1))

lstm_model.compile(optimizer=Adam(learning_rate=0.001), loss='mse')

# Train (example for LSTM)

early_stop = EarlyStopping(monitor='val_loss', patience=100)

history_lstm = lstm_model.fit(X_train, y_train, epochs=100, batch_size=32,
validation_split=0.2, callbacks=[early_stop])

early_stop = EarlyStopping(monitor='val_loss', patience=100)

history_rnn = rnn_model.fit(X_train, y_train, epochs=100, batch_size=32,
validation_split=0.2, callbacks=[early_stop])

# Save models

rnn_model.save('rnn_model.keras')

lstm_model.save('lstm_model.keras')
```

## test.py code

```
import pandas as pd

from tensorflow.keras.models import load_model

from sklearn.metrics import mean_absolute_error, mean_squared_error

import numpy as np

import matplotlib.pyplot as plt

import pickle

rnn_model = load_model('rnn_model.keras')

lstm_model = load_model('lstm_model.keras')

with open('scaler.pkl', 'rb') as f:

    scaler = pickle.load(f)

X_test = np.load('X_test.npy')

y_test = np.load('y_test.npy')

test_data = pd.read_pickle('test_data.pkl')

with open('window_size.txt', 'r') as f:

    window_size = int(f.read())

# Predict

rnn_preds = rnn_model.predict(X_test)

lstm_preds = lstm_model.predict(X_test)

# Inverse scale

rnn_preds = scaler.inverse_transform(rnn_preds)
```

```
lstm_preds = scaler.inverse_transform(lstm_preds)

y_test_inv = scaler.inverse_transform(y_test.reshape(-1, 1))

# Metrics for both models

rnn_mae = mean_absolute_error(y_test_inv, rnn_preds)

rnn_rmse = np.sqrt(mean_squared_error(y_test_inv, rnn_preds))

lstm_mae = mean_absolute_error(y_test_inv, lstm_preds)

lstm_rmse = np.sqrt(mean_squared_error(y_test_inv, lstm_preds))

print(f'RNN: MAE={rnn_mae:.2f}°C, RMSE={rnn_rmse:.2f}°C')
print(f'LSTM: MAE={lstm_mae:.2f}°C, RMSE={lstm_rmse:.2f}°C')

# Plot Actual vs Predicted for RNN

plt.figure(figsize=(12, 6))

plt.plot(test_data.index[window_size:], y_test_inv, label='Actual')

plt.plot(test_data.index[window_size:], rnn_preds, label='RNN Predicted')

plt.legend()

plt.title('RNN: Actual vs Predicted (2025 Test Data)')

plt.xlabel('Date')

plt.ylabel('Temperature (°C)')

plt.savefig('rnn_actual_vs_predicted.png') # Save for report/GitHub

plt.show()
```

```
# Plot Actual vs Predicted for LSTM

plt.figure(figsize=(12, 6))

plt.plot(test_data.index[window_size:], y_test_inv, label='Actual')

plt.plot(test_data.index[window_size:], lstm_preds, label='LSTM Predicted')

plt.legend()

plt.title('LSTM: Actual vs Predicted (2025 Test Data)')

plt.xlabel('Date')

plt.ylabel('Temperature (°C)')

plt.savefig('lstm_actual_vs_predicted.png') # Save for report/GitHub

plt.show()
```

```
# Error distribution for LSTM (example; repeat for RNN if needed)
```

```
errors = y_test_inv - lstm_preds

plt.hist(errors, bins=30)

plt.title('LSTM Error Distribution')

plt.xlabel('Error (°C)')

plt.ylabel('Frequency')

plt.savefig('lstm_error_distribution.png')

plt.show()
```

```
# Error distribution for RNN (example; repeat for RNN if needed)
```

```
errors = y_test_inv - rnn_preds

plt.hist(errors, bins=30)

plt.title('RNN Error Distribution')

plt.xlabel('Error (°C)')
```

```

plt.ylabel('Frequency')

plt.savefig('rnn_error_distribution.png')

plt.show()

# Optional: Confidence intervals (simple example using rolling std dev on
# predictions)

rolling_std = pd.Series(lstm_preds.flatten()).rolling(window=7).std().values
# 7-day rolling std

plt.figure(figsize=(12, 6))

plt.plot(test_data.index[window_size:], y_test_inv, label='Actual')

plt.plot(test_data.index[window_size:], lstm_preds, label='LSTM Predicted')

plt.fill_between(test_data.index[window_size:],      lstm_preds.flatten() -
rolling_std,  lstm_preds.flatten() + rolling_std,  color='gray', alpha=0.2,
label='Confidence Interval (±1 std)')

plt.legend()



plt.title('LSTM with Confidence Intervals (2025 Test Data)')

plt.savefig('lstm_with_ci.png')


plt.show()

```


## Models predicted result




LSTM Error Distribution



RNN Error Distribution



## Models testing



```
命令提示字元 - test.py
Microsoft Windows [版本 10.0.26200.7309]
(c) Microsoft Corporation. 著作權所有，並保留一切權利。
C:\Users\Zac>cd C:\Users\Zac\Downloads\5411Project
C:\Users\Zac\Downloads\5411Project>test.py
2025-12-08 23:22:57.111700: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable "TF_ENABLE_ONEDNN_OPTS=0".
2025-12-08 23:23:01.240034: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable "TF_ENABLE_ONEDNN_OPTS=0".
2025-12-08 23:23:04.316709: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: SSE3 SSE4.1 SSE4.2 AVX AVX2 AVX_VNNI FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
[[1m9/90[0m[32m-----+ [0m[37m[0m[[1m0s\0m 14ms/step
[[1m9/90[0m[32m-----+ [0m[37m[0m[[1m0s\0m 14ms/step
RNN: MAE=0.96°C, RMSE=1.31°C
LSTM: MAE=1.10°C, RMSE=1.38°C
```

# Error Analysis: Impact of Extreme Weather Events on Model Predictions

A detailed examination of prediction errors reveals that deviations are particularly pronounced during extreme weather events, such as black rainstorms and heavy rainfall in August 2025. In Hong Kong's subtropical climate, black rainstorms—defined as rainfall exceeding 70 mm/hour—trigger ground cooling through evaporation and convection, leading to unexpected drops in minimum temperatures despite typically stable summer highs (around 25–27°C). 2025 recorded a historic 4 black rainstorm warnings, the highest since records began in 1884, with three occurring in early August.

## 2025年8月初的連場黑雨

星期二, 2025年8月5日

今年香港迎來一個多雨的夏天，於7月29日出現今年首場黑雨，雖然隨後一兩日天氣稍為好轉，但自上週六（8月2日）開始天氣持續不穩定，發生連場暴雨，天文台需於4日內三次發出黑色暴雨警告信號。連同7月29日的暴雨，今年至今已出現四次黑雨，打破一年內黑雨生效次數的紀錄。



## References

1. Hong Kong Observatory. (2025). *Daily Grass Minimum Temperature Dataset*. data.gov.hk. <https://data.gov.hk/en-data/dataset/hk-hko-rss-daily-grass-min-temp>